

# Learning in Intelligent Systems

Artificial Intelligence @ Allegheny College

Janyl Jumadinova

September 13-17, 2021

# Overview of Learning



# Learning in Humans



- The act / process of acquiring, modify or reinforcing knowledge or skills through synthesizing different types of new or existed information.
- Key to human survival.

# Learning in Humans



- The act / process of acquiring, modify or reinforcing knowledge or skills through synthesizing different types of new or existed information.
- Key to human survival.
- Progress over time tends to follow learning curves (relatively permanent).

# Learning in Computing Systems



- Computational methods using “experience” to improve performance.

# Learning in Computing Systems



- Computational methods using “experience” to improve performance.
- Experience — data driven task.

# Learning in Computing Systems



- Computational methods using “experience” to improve performance.
- Experience — data driven task.
- Computer science — involves learning algorithms, analysis of complexity, and theoretical guarantees.



# Learning in Computing Systems

Artificial intelligence | Machine learning

# Learning in Computing Systems

## Artificial intelligence | Machine learning

- Computer program(s) with adaptive mechanisms that enable computer / machine to learn from experience / example / analogy / rewards.

# Learning in Computing Systems

## Artificial intelligence | Machine learning

- Computer program(s) with adaptive mechanisms that enable computer / machine to learn from experience / example / analogy / rewards.
- It improves the performance of an intelligent system over time (e.g, reducing error rate, improving rewards).

# Why Learning in Computing Systems?

- Understand and improve efficiency of human learning / understanding.

# Why Learning in Computing Systems?

- Understand and improve efficiency of human learning / understanding.
- Discover new things or structure that is unknown to humans.

# Why Learning in Computing Systems?

- Understand and improve efficiency of human learning / understanding.
- Discover new things or structure that is unknown to humans.
- Fill in skeletal or incomplete knowledge / expert specifications about a domain.

# Applications of Learning

Mainly in decision making / pattern recognition / **intelligent systems**.


# Applications of Learning

Mainly in decision making / pattern recognition / **intelligent systems**.


- Robot navigation.
- Automatic speech recognition (Siri in iPhone, Google speech-to-text search).
- Search and recommendation (Google, Amazon, eBay).
- Financial prediction, fraud detection, medical diagnosis.
- Video games, data visualization.



# Black-box Learning



# Learning Architecture



# Learning Paradigms

- **Supervised learning**
  - input-output relationships

# Learning Paradigms

- **Supervised learning**
  - input-output relationships
- **Unsupervised learning**
  - relationship among inputs

# Learning Paradigms

- **Supervised learning**
  - input-output relationships
- **Unsupervised learning**
  - relationship among inputs
- **Reinforcement learning**
  - input-action relates to rewards / punishment

# Supervised Learning

Given examples of inputs and corresponding desired outputs.

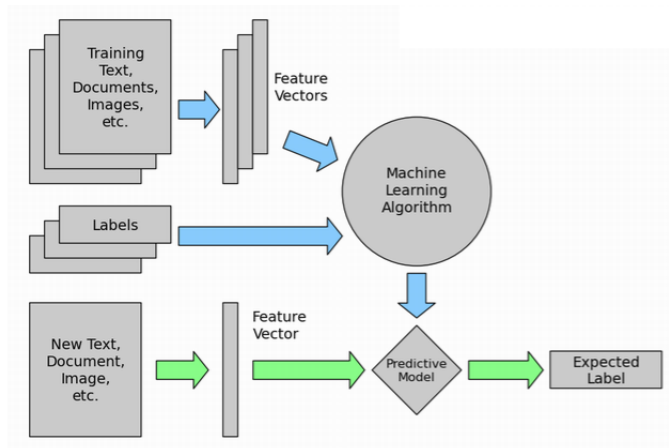
# Supervised Learning

Given examples of inputs and corresponding desired outputs.

## Tasks:

- **Classification** (categorizing output: correct class)
- **Regression** (continuous output to predict output based for new inputs)
- **Prediction** (classify / regression on new input sequences)

# Supervised Learning





# Unsupervised Learning

Given only inputs and automatically discover representations, features, structure etc.

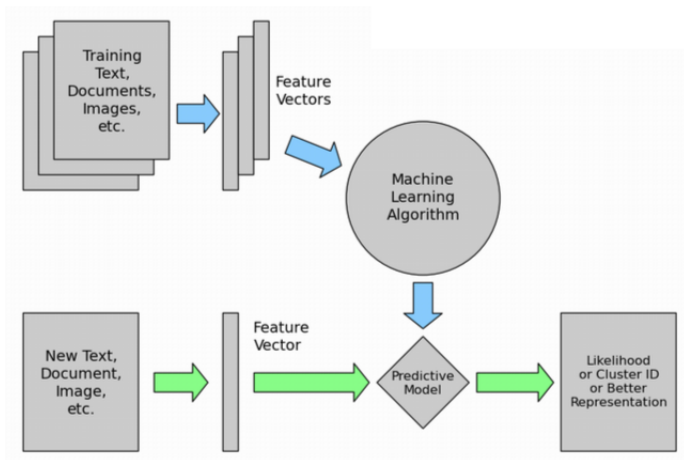
# Unsupervised Learning

Given only inputs and automatically discover representations, features, structure etc.

## Tasks:

- **Clustering** (to group similar data into a finite number of clusters / groups)
- **Vector Quantization** (compress / decode dataset into a new representation but maintaining internal information)
- **Outlier Detection** (select highly unusual cases/sequences)

# Unsupervised Learning



# Reinforcement Learning

- Learning approach of getting a computer system to act in the world so as to maximize its rewards.


# Reinforcement Learning

- Learning approach of getting a computer system to act in the world so as to maximize its rewards.
- Consider teaching a domestic animal. We cannot tell it what to do, but we can reward / punish if it does the right/ wrong thing.

# Reinforcement Learning

- Learning approach of getting a computer system to act in the world so as to maximize its rewards.
- Consider teaching a domestic animal. We cannot tell it what to do, but we can reward / punish if it does the right/ wrong thing.
- Process to determine what it did that made it get the reward / punishment – “credit assignment problem.”

# Reinforcement Learning



# Learning Lifecycle



<https://www.openshift.com/>



## Activity 5: Algorithmic Bias

# Google Search

# Supervised Learning

# Supervised Learning: Performance Measures

- A **feature** is a measurable property or a characteristic of the object we are trying to analyze (columns in a data set).


# Supervised Learning: Performance Measures

- A **feature** is a measurable property or a characteristic of the object we are trying to analyze (columns in a data set).
- **Discrimination** attempts to separate distinct sets of objects.
- **Classification** attempts to allocate new objects to predefined groups.


# Performance Measures

- **Cost ratio** is a ratio of *false positives* (given condition is present when it is not) to *false negatives* (given condition is not present when it actually is).
- **Confusion matrix** (error matrix): a table to visualize the performance of an algorithm with rows/columns representing instances of predictions and columns/rows representing instances of actual class.

# Confusion Matrix



# Confusion Matrix



- a is a *true positive (TP)*.
- d is a *true negative (TN)*.
- c is a *false positive (FP)*.
- b is a *false negative (FN)*.

# Classification Accuracy

Number of correctly classified examples divided by the total number of examples.

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$



# Classification Accuracy

Number of correctly classified examples divided by the total number of examples.

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (1)$$

$$Error = 1 - Accuracy \quad (2)$$

# Performance Measures

$$Recall = \frac{TP}{TP + FN} \quad (3)$$

Higher the recall the better class is correctly recognized (small number of FN).

# Performance Measures

$$Recall = \frac{TP}{TP + FN} \quad (3)$$

Higher the recall the better class is correctly recognized (small number of FN).

$$Precision = \frac{TP}{TP + FP} \quad (4)$$

Higher the precision the better indication of an example labeled as positive being indeed positive (small number of FP).

# Performance Measures

- High recall, low precision: Most of the positive examples are correctly recognized (low FN) but there are a lot of false positives.
- Low recall, high precision: Miss a lot of positive examples (high FN) but those we predict as positive are indeed positive (low FP).

# Performance Measures

- High recall, low precision: Most of the positive examples are correctly recognized (low FN) but there are a lot of false positives.
- Low recall, high precision: Miss a lot of positive examples (high FN) but those we predict as positive are indeed positive (low FP).

$$F1 = 2 \frac{Precision * Recall}{Precision + Recall} \quad (5)$$

F1 Score is used to find a balance between Precision and Recall.

# Performance Measures

- Receiver Operator Characteristic **ROC curve**: plot of TP vs. FP.

# Performance Measures

- Receiver Operator Characteristic **ROC curve**: plot of TP vs. FP.
- ROC Curves summarize the trade-off between the true positive rate and false positive rate for a predictive model using different probability thresholds.

# Performance Measures

- Receiver Operator Characteristic **ROC curve**: plot of TP vs. FP.
- ROC Curves summarize the trade-off between the true positive rate and false positive rate for a predictive model using different probability thresholds.
- Precision-Recall curves summarize the trade-off between the true positive rate and the positive predictive value for a predictive model using different probability thresholds.



# Performance Measures

- Receiver Operator Characteristic **ROC curve**: plot of TP vs. FP.
- ROC Curves summarize the trade-off between the true positive rate and false positive rate for a predictive model using different probability thresholds.
- Precision-Recall curves summarize the trade-off between the true positive rate and the positive predictive value for a predictive model using different probability thresholds.
- ROC curves are appropriate when the observations are balanced between each class, whereas precision-recall curves are appropriate for imbalanced datasets.

# What-If Tool

## Smile Detection Demo

<https://pair-code.github.io/what-if-tool/>

# Computer Vision

Make computers understand images and video.

# Computer Vision

Make computers understand images and video.



# Computer Vision

Make computers understand images and video

# Computer Vision

Make computers understand images and video




# Computer Vision

- What kind of scene?
- Where are the cars?
- How far is the building?

# Why computer vision matters?



Safety



Health




Security



Comfort



Fun




Access



## Applications of Computer Vision



"Face Recognition"




"Pose Estimation"



"Body Tracking"



"Speech Reading"



"Palm Recognition"



"Car Tracking"

# Segmentation

- Compact representation for image data in terms of a set of **components**.

# Segmentation


- Compact representation for image data in terms of a set of **components**.
- Components share “common” visual **properties**.

# Segmentation

- Compact representation for image data in terms of a set of **components**.
- Components share “common” visual **properties**.
- Properties can be defined at different level of **abstractions**.

# Segmentation

- Compact representation for image data in terms of a set of **components**.
- Components share “common” visual **properties**.
- Properties can be defined at different level of **abstractions**.



From: <https://docs.opencv.org>

# Segmentation

- Tokens
  - whatever we need to group (pixels, points, surface elements, etc.).

# Segmentation

- Tokens
  - whatever we need to group (pixels, points, surface elements, etc.).
- Bottom up segmentation
  - tokens belong together because they are locally coherent.

# Segmentation

- Tokens
  - whatever we need to group (pixels, points, surface elements, etc.).
- Bottom up segmentation
  - tokens belong together because they are locally coherent.
- Top down segmentation
  - tokens belong together because they lie on the same object.



# Segmentation

- **Tokens**
  - whatever we need to group (pixels, points, surface elements, etc.).
- **Bottom up segmentation**
  - tokens belong together because they are locally coherent.
- **Top down segmentation**
  - tokens belong together because they lie on the same object.




Figure 1: The relative merits of the bottom-up and the top-down

# What is Segmentation?


Clustering image elements that “belong together”

- **Partitioning**
  - Divide into regions/sequences with coherent internal properties.
- **Grouping**
  - Identify sets of coherent tokens in image.

# OpenCV

- An open source BSD licensed computer vision library
  - Patent-encumbered code isolated into “non-free” module (SIFT, SURF, some of the Face Detectors, etc.)
- Available on all major platforms
  - Android, iOS, Linux, Mac OS X, Windows
- Written primarily in C++
  - Bindings available for Python, Java, even MATLAB (in 3.0).
- Well documented at <http://docs.opencv.org>
- Source available at <https://github.com/Itseez/opencv>

# OpenCV



# OpenCV: Pixel

- **Grayscale:** each pixel has a value between 0 (black) and 255 (white)
  - values between 0 and 255 are varying shades of gray.

# OpenCV: Pixel

- **Grayscale:** each pixel has a value between 0 (black) and 255 (white)
  - values between 0 and 255 are varying shades of gray.
- **Color:** pixels are normally represented in the RGB color space
  - one value for the Red component, one for Green, and one for Blue,
  - each of the three colors is represented by an integer in the range 0 to 255,
  - how “much” of the color there is.

# OpenCV: Coordinate System

- The point  $(0, 0)$  corresponds to the upper left corner of the image
- x value increases as we move to the right
- y value increases as we move down

# OpenCV: Image Representation

- OpenCV represents images as NumPy arrays (matrices).
- NumPy is a library for the Python programming language that provides support for large, multi-dimensional arrays.
- To access a pixel value, we need to supply the  $x$  and  $y$  coordinates of the pixel.
- OpenCV actually stores RGB values in the order of Blue, Green, and Red.



# Images

How to input or output an image?

## How to input or output an image?

### Load Image

*Read image from disk.*

```
cv::imread(filename, 0/1);
```

0: read as grayscale image

1: read as color image

### Save Image

*Write image to disk.*

```
cv::imwrite(filename, im);
```

### Visualize Image

*Show image in a window.*

```
cv::imshow(title, im);
```

*Note: if CV\_32FC1, the gray value range is 0 to 1. Everything above 1 is white and everything below 0 is black.*

### Waitkey


*Waits n milliseconds for user input.*

```
cv::waitkey(n);
```


*If n == -1, it waits forever.*

*Note: There must be a waitkey to show the image.*

# Drawing Primitives



```
cv::line(im, p1, p2, color, thickness);  
           cv::Point(x, y)
```



```
cv::circle(im, c, r, color, thickness);  
            CV_RGB(r, g, b)
```

# Drawing Primitives

```
rectangle = np.zeros((300, 300), dtype = "uint8")  
cv2.rectangle(rectangle, (25, 25), (275, 275), 255, -1)
```

# Bitwise Operations

Examine every pixel in the input images:

- `cv2.bitwise_and` (used in masking example): if both pixels have a value  $> 0$ , the output pixel is set to 255 in the output image, otherwise it is 0.
- `cv2.bitwise_or`: if either of the pixels have a value  $> 0$ , the output pixel is set to 255 in the output image, otherwise it is 0.
- `cv2.bitwise_xor`: same as OR, with a restriction: both pixels are not allowed to have values  $> 0$ .
- `cv2.bitwise_not`: pixels with a value of 255 become 0, pixels with a value of 0 become 255.

- ① Load an image from the disk, display it on our screen, and write it to file in a different format.
- ② Access and manipulate pixels.