

# Summary/Planning

Artificial Intelligence @ Allegheny College

Janyl Jumadinova

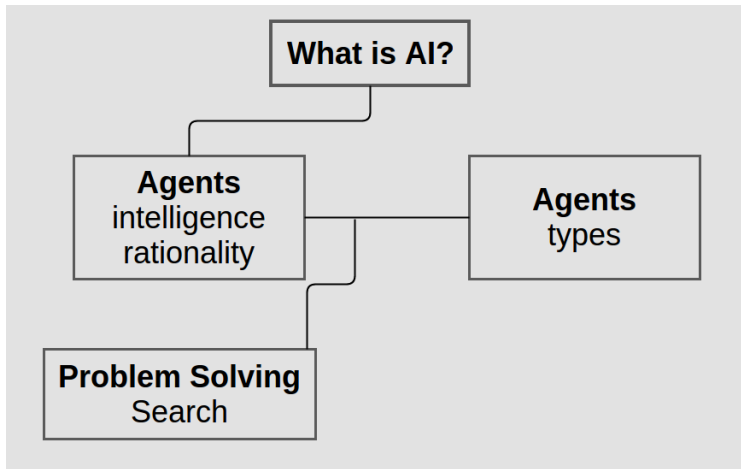
December 1, 2021

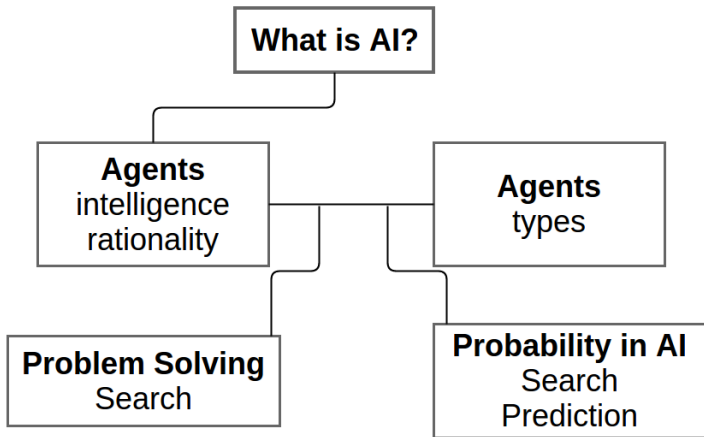
# What is AI?

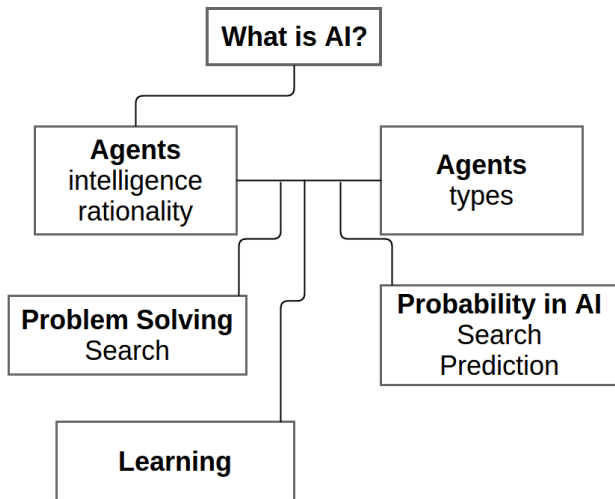
**What is AI?**

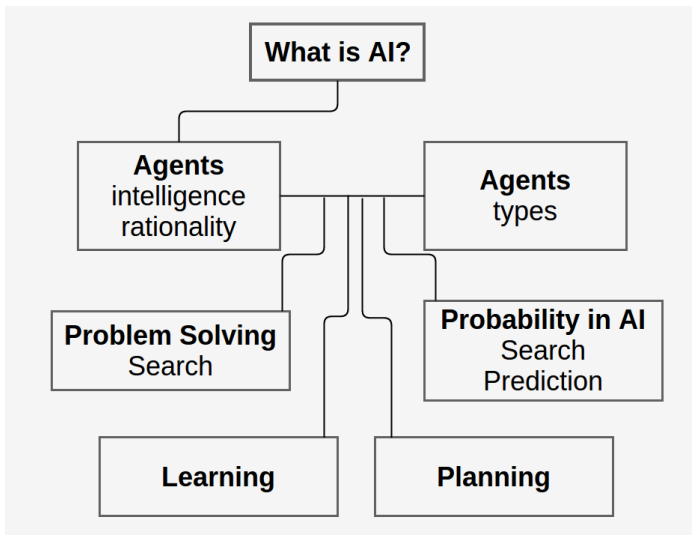
**Agents**  
intelligence  
rationality













# Planning

## Planning:

- deciding what to do based on an agent's ability, its goals, and the state of the world;
- 'thinking head' using logical representations of the states of the world;
- finding a sequence of actions to solve a goal.

# Planning

## Planning:

- deciding what to do based on an agent's ability, its goals, and the state of the world;
- 'thinking head' using logical representations of the states of the world;
- finding a sequence of actions to solve a goal.

In *classical planning*, the environment is fully observable; deterministic and static with only a single agent.

# Planning

- Automation requires efficient automated planning.
- In comparison with the classification problem planning problem solutions provide guarantees on its quality.

# Planning

- Automation requires efficient automated planning.
- In comparison with the classification problem planning problem solutions provide guarantees on its quality.

## **When to use Planning:**

- Desire a procedural course of action for a declaratively described system.
- Domain knowledge can be elicited or learned over time.
- Consistency is more desired than learning transient behaviors.


# Applications

- System control (autonomous/virtual agents)
- Process control:
  - Construction and configuration
  - Workflow management
  - Mission planning
  - Project planning

Basic planning problem:

- **Given:** start state, goal conditions, actions
- **Find:** sequence of actions leading from start to goal
- **Typically:** states correspond to possible worlds; actions and goals specified using a logical formalism (e.g., STRIPS, situation calculus, temporal logic, etc.)

# Delivery Robot Example



## Features:

*RLoc* – Rob's location

*RHC* – Rob has coffee

*SWC* – Sam wants coffee

*MW* – Mail is waiting

*RHM* – Rob has mail

## Actions:

*mc* – move clockwise

*mcc* – move counterclockwise

*puc* – pickup coffee

*dc* – deliver coffee

*pum* – pickup mail

*dm* – deliver mail

# Explicit State-space Representation

Enumerate the states and, for each state, specify the actions that are possible in that state and, for each state-action pair, specify the state that results from carrying out the action in that state.



# Explicit State-space Representation

Enumerate the states and, for each state, specify the actions that are possible in that state and, for each state-action pair, specify the state that results from carrying out the action in that state.

$$4 \times 2 \times 2 \times 2 \times 2 = 64 \text{ states}$$

State	Action	Resulting State
$\langle lab, \overline{rhc}, swc, \overline{mw}, rhm \rangle$	<i>mc</i>	$\langle mr, \overline{rhc}, swc, \overline{mw}, rhm \rangle$
$\langle lab, \overline{rhc}, swc, \overline{mw}, rhm \rangle$	<i>mcc</i>	$\langle off, \overline{rhc}, swc, \overline{mw}, rhm \rangle$
$\langle off, \overline{rhc}, swc, \overline{mw}, rhm \rangle$	<i>dm</i>	$\langle off, \overline{rhc}, \overline{swc}, \overline{mw}, rhm \rangle$
$\langle off, \overline{rhc}, swc, \overline{mw}, rhm \rangle$	<i>mcc</i>	$\langle cs, \overline{rhc}, swc, \overline{mw}, rhm \rangle$
$\langle off, \overline{rhc}, swc, \overline{mw}, rhm \rangle$	<i>mc</i>	$\langle lab, \overline{rhc}, swc, \overline{mw}, rhm \rangle$
...	...	...

# Status of Classical Planning

- **Classical planning works!**
  - Large problems solved very fast (non-optimally)

# Status of Classical Planning

- **Classical planning works!**
  - Large problems solved very fast (non-optimally)
- **Limitations:**
  - Does not model **uncertainty** (no probabilities)
  - Does not deal with **incomplete information** (no sensing)
  - Deals with very **simple cost structure** (no state dependent costs)

Static vs. Dynamic

*Environment*

Fully  
vs.  
Partially  
Observable

Deterministic  
vs.  
Stochastic

Sequential  
vs.  
Concurrent

Instantaneous  
vs.  
Durative

Perfect  
vs.  
Noisy

*Percepts*

*Actions*



# Policies

A **stationary policy** is a function:


$$\pi : S \rightarrow A$$

Given a state  $s$ ,  $\pi(s)$  specifies what action the agent who is following  $\pi$  will do.

An **optimal policy** is one with maximum expected discounted reward.

# Policy Examples

- ① Never check for smoke, and call only if there is a report.
- ② Always check for smoke, and call only if it sees smoke.
- ③ Check for smoke if there is a report, and call only if there is a report and it sees smoke.
- ④ Check for smoke if there is no report, and call when it does not see smoke.
- ⑤ Always check for smoke and never call.



# Markov Decision Processes (MDPs)

A fundamental framework for probabilistic planning.

**MDPs** are fully observable, probabilistic state models:

- a state space  $S$
- a set  $G \subset S$  of goal states
- actions  $A(s) \subset A$  applicable in each state  $s \in S$
- transition probabilities  $P_a(s_0|s)$  for  $s \in S$  and  $a \in A(s)$
- action costs  $c(a, s) > 0$

# Markov Decision Processes (MDPs)

A fundamental framework for probabilistic planning.

**MDPs** are fully observable, probabilistic state models:


- a state space  $S$
- a set  $G \subset S$  of goal states
- actions  $A(s) \subset A$  applicable in each state  $s \in S$
- transition probabilities  $P_a(s_0|s)$  for  $s \in S$  and  $a \in A(s)$
- action costs  $c(a, s) > 0$

**Solutions** are functions (**policies**) mapping states into actions.

**Optimal** solutions have minimum **expected** costs.



# Markov Decision Processes (MDPs)



# MDP Solution


- Want a way to choose an action in a state, i.e., a policy  $\pi$
- What does a policy look like?
  - Can pick action based on states visited + actions used so far, i.e., execution history  $h = s(1)a(1)s(2)a(2)\dots$
  - Can pick actions randomly

# MDP Solution

- Want a way to choose an action in a state, i.e., a policy  $\pi$
- What does a policy look like?
  - Can pick action based on states visited + actions used so far, i.e., execution history  $h = s(1)a(1)s(2)a(2)\dots$
  - Can pick actions randomly
- Thus, in general an MDP solution is a probabilistic history-dependent  $\pi : H \times A \rightarrow [0, 1]$

# Evaluating MDP Solutions

- Executing a policy yields a sequence of rewards



# Evaluating MDP Solutions

- Define utility function  $u(R1, R2, \dots)$  to be some “quality measure” of a reward sequence.
- Define **value** function as  $V : H \rightarrow [-\infty, \infty]$ .
- Define value function of a policy after history  $h$  to be some utility function of subsequent rewards:  
$$V^\pi(h) = u(R1, R2, \dots)$$

# Optimal MDP Solution

- **Want:** a behavior that is “best” in every situation
- $\pi^*$  is an optimal policy if  $V^*(h) \geq V^\pi(h)$  for all  $\pi$ , for all  $h$

Example: to exercise or not?

Example: to exercise or not?

Each week *Sam* has to decide whether to exercise or not:

**States:** {fit, unfit} **Actions:** {exercise, relax}



Example: to exercise or not?

Each week *Sam* has to decide whether to exercise or not:

**States:** {fit, unfit} **Actions:** {exercise, relax}

Dynamics:

State	Action	$P(\text{fit}   \text{State}, \text{Action})$
fit	exercise	0.99
fit	relax	0.7
unfit	exercise	0.2
unfit	relax	0.0

Reward (does not depend on resulting state):

State	Action	Reward
fit	exercise	8
fit	relax	10
unfit	exercise	0
unfit	relax	5

# Information Availability

What information is available when the agent decides what to do?

- **Fully-observable MDP** the agent gets to observe  $S_t$  when deciding on action  $A_t$ .

# Information Availability

What information is available when the agent decides what to do?

- **Fully-observable MDP** the agent gets to observe  $S_t$  when deciding on action  $A_t$ .
- **Partially-observable MDP (POMDP)** the agent has some noisy sensor of the state. It needs to remember its sensing and acting history.

# AI Planning at IBM

