

Basics of Probability

Basics of Probability

Janyl Jumadinova

October 4–11, 2021

Probability Theory

- **Probability theory yields mathematical tools to deal with uncertain events.**

Probability Theory

- Probability theory yields mathematical tools to deal with uncertain events.
- **Used everywhere nowadays and its importance is growing.**

Probability and Statistics

- Probability \neq Statistics
- Probability: Known distributions \Rightarrow what are the outcomes?
- Statistics: Known outcomes \Rightarrow what are the distributions?

Counting

- Many **basic** probability problems are counting problems.

Counting

- Many **basic** probability problems are counting problems.
- *Example:* Assume there are 1 man and 2 women in a room. You pick a person randomly. What is the probability P_1 that this is a man?

Counting

- Many **basic** probability problems are counting problems.
- *Example:* Assume there are 1 man and 2 women in a room. You pick a person randomly. What is the probability P_1 that this is a man?

Counting

- Many **basic** probability problems are counting problems.
- *Example:* Assume there are 1 man and 2 women in a room. You pick a person randomly. What is the probability P_1 that this is a man? If you pick two persons randomly, what is the probability P_2 that these are a man and woman?
- *Answer:* You have the possible outcomes: (M), (W1), (W2) so

$$P_1 = \frac{\# \text{ "successful" events}}{\# \text{ events}} = \frac{\# \text{ men}}{\# \text{ men} + \# \text{ women}} = \frac{1}{3}.$$

To compute P_2 , you can think of all the possible events: (M,W1), (M,W2), (W1,W2) so

$$P_2 = \frac{\# \text{ "successful" events}}{\# \text{ events}} = \frac{2}{3}.$$

Sample Space

Definition

The *sample space* S of an experiment (whose outcome is uncertain) is the set of all possible outcomes of the experiment.

Sample Space

- *Example* (**child**): Determining the sex of a newborn child in which case $S = \{boy, girl\}$.

Sample Space

- *Example* (child): Determining the sex of a newborn child in which case $S = \{boy, girl\}$.
- *Example* (horse race): Assume you have an horse race with 12 horses. If the experiment is the order of finish in a race, then

$$S = \{\text{all } 12! \text{ permutations of } (1, 2, 3, \dots, 11, 12)\}.$$

Sample Space

- *Example* (**child**): Determining the sex of a newborn child in which case $S = \{\text{boy}, \text{girl}\}$.
- *Example* (**horse race**): Assume you have an horse race with 12 horses. If the experiment is the order of finish in a race, then

$$S = \{\text{all } 12! \text{ permutations of } (1, 2, 3, \dots, 11, 12)\}.$$

- *Example* (**coins**): If the experiment consists of flipping two coins, then the sample space is

$$S = \{(H, H), (H, T), (T, H), (T, T)\}.$$

Sample Space

- *Example (child)*: Determining the sex of a newborn child in which case $S = \{\text{boy}, \text{girl}\}$.
- *Example (horse race)*: Assume you have an horse race with 12 horses. If the experiment is the order of finish in a race, then

$$S = \{\text{all } 12! \text{ permutations of } (1, 2, 3, \dots, 11, 12)\}.$$

- *Example (coins)*: If the experiment consists of flipping two coins, then the sample space is

$$S = \{(H, H), (H, T), (T, H), (T, T)\}.$$

- *Example (lifetime)*: If the experiment consists of measuring the lifetime (in years) of your pet then the sample space consists of all nonnegative real numbers: $S = \{x; 0 \leq x < \infty\}$.

Events

- Any *subset* E of the sample space S is known as an *event*; i.e. an event is a set consisting of possible outcomes of the experiment.

Events

- Any *subset* E of the sample space S is known as an *event*; i.e. an event is a set consisting of possible outcomes of the experiment.
- If the outcome of the experiment is in E , then we say that E has occurred.

Events

- *Example* (child): The event $E = \{boy\}$ is the event that the child is a boy.

Events

- *Example* (child): The event $E = \{boy\}$ is the event that the child is a boy.
- *Example* (horse race): The event $E = \{\text{all outcomes in } S \text{ starting with a 7}\}$ is the event that the race was won by horse 7.

Events

- *Example* (**child**): The event $E = \{\text{boy}\}$ is the event that the child is a boy.
- *Example* (**horse race**): The event $E = \{\text{all outcomes in } S \text{ starting with a } 7\}$ is the event that the race was won by horse 7.
- *Example* (**coins**): The event $E = \{(H, T), (T, T)\}$ is the event that a tail appears on the second coin.

Events

- *Example* (**child**): The event $E = \{\text{boy}\}$ is the event that the child is a boy.
- *Example* (**horse race**): The event $E = \{\text{all outcomes in } S \text{ starting with a 7}\}$ is the event that the race was won by horse 7.
- *Example* (**coins**): The event $E = \{(H, T), (T, T)\}$ is the event that a tail appears on the second coin.
- *Example* (**lifetime**): The event $E = \{x : 3 \leq x \leq 15\}$ is the event that your pet will live more than 3 years but won't live more than 15 years.

Union of Events

Given events E and F , $E \cup F$ is the set of all outcomes *either* in E or F or in *both* E and F .

$E \cup F$ occurs if *either* E or F occurs.

$E \cup F$ is the **union** of events E and F

Union of Events

- *Example (coins)*: If we have $E = \{(H, T)\}$ and $F = \{(T, H)\}$ then $E \cup F = \{(H, T), (T, H)\}$ is the event that one coin is head and the other is tail.

Union of Events

- *Example (coins)*: If we have $E = \{(H, T)\}$ and $F = \{(T, H)\}$ then $E \cup F = \{(H, T), (T, H)\}$ is the event that one coin is head and the other is tail.
- *Example (horse race)*: If we have $E = \{\text{all outcomes in } S \text{ starting with a } 7\}$ and $F = \{\text{all outcomes in } S \text{ finishing with a } 3\}$ then $E \cup F$ is the event that the race was won by horse 7 and/or the last horse was horse 3.

Union of Events

- *Example (coins)*: If we have $E = \{(H, T)\}$ and $F = \{(T, H)\}$ then $E \cup F = \{(H, T), (T, H)\}$ is the event that one coin is head and the other is tail.
- *Example (horse race)*: If we have $E = \{\text{all outcomes in } S \text{ starting with a } 7\}$ and $F = \{\text{all outcomes in } S \text{ finishing with a } 3\}$ then $E \cup F$ is the event that the race was won by horse 7 and/or the last horse was horse 3.
- *Example (lifetime)*: If $E = \{x : 0 \leq x \leq 10\}$ and $F = \{x : 15 \leq x < \infty\}$ then $E \cup F$ is the event that your pet will die before 10 or will die after 15.

Intersection of Events

Given events E and F , $E \cap F$ is the set of all outcomes which are *both* in E and F .

$E \cap F$ is also denoted as EF .

Intersection of Events

- *Example (coins):* If we have $E = \{(H, H), (H, T), (T, H)\}$ (event that one H at least occurs) and $F = \{(H, T), (T, H), (T, T)\}$ (event that one T at least occurs) then $E \cap F = \{(H, T), (T, H)\}$ is the event that one H and one T occur.

Intersection of Events

- *Example (coins)*: If we have $E = \{(H, H), (H, T), (T, H)\}$ (event that one H at least occurs) and $F = \{(H, T), (T, H), (T, T)\}$ (event that one T at least occurs) then $E \cap F = \{(H, T), (T, H)\}$ is the event that one H and one T occur.
- *Example (horse race)*: If we have $E = \{\text{all outcomes in } S \text{ starting with a 7}\}$ and $F = \{\text{all outcomes in } S \text{ starting with a 8}\}$ then $E \cap F$ does not contain any outcome and is denoted by \emptyset .

Intersection of Events

- *Example (coins)*: If we have $E = \{(H, H), (H, T), (T, H)\}$ (event that one H at least occurs) and $F = \{(H, T), (T, H), (T, T)\}$ (event that one T at least occurs) then $E \cap F = \{(H, T), (T, H)\}$ is the event that one H and one T occur.
- *Example (horse race)*: If we have $E = \{\text{all outcomes in } S \text{ starting with a 7}\}$ and $F = \{\text{all outcomes in } S \text{ starting with a 8}\}$ then $E \cap F$ does not contain any outcome and is denoted by \emptyset .
- *Example (lifetime)*: If we have $E = \{x : 0 \leq x \leq 5\}$ and $F = \{x : 10 \leq x < 15\}$ then $E \cap F = \{x : 3 \leq x \leq 5\}$ is the event that your pet will die between 10 and 15.

Notations and Properties

- For any event E , E^c denote the *complement* set of all outcomes in S which are not in E .

Hence we have $E \cup E^c = S$ and $E \cap E^c = \emptyset$.

Notations and Properties

- For any event E , E^c denote the *complement* set of all outcomes in S which are not in E .

Hence we have $E \cup E^c = S$ and $E \cap E^c = \emptyset$.

- For any two events E and F , we write $E \subset F$ if all the outcomes of E are in F .

Axioms of Probability

- Consider an experiment with sample space S . For each event E , we assume that a number $P(E)$, the *probability* of the event E , is defined and satisfies the following 3 axioms.

Axioms of Probability

- Consider an experiment with sample space S . For each event E , we assume that a number $P(E)$, the *probability* of the event E , is defined and satisfies the following 3 axioms.
- Axiom 1**

$$0 \leq P(E) \leq 1$$

Axioms of Probability

- Consider an experiment with sample space S . For each event E , we assume that a number $P(E)$, the *probability* of the event E , is defined and satisfies the following 3 axioms.

- Axiom 1**

$$0 \leq P(E) \leq 1$$

- Axiom 2**

$$P(S) = 1$$

Axioms of Probability

- Consider an experiment with sample space S . For each event E , we assume that a number $P(E)$, the *probability* of the event E , is defined and satisfies the following 3 axioms.

- Axiom 1**

$$0 \leq P(E) \leq 1$$

- Axiom 2**

$$P(S) = 1$$

- Axiom 3.** For any sequence of mutually exclusive events $\{E_i\}_{i \geq 1}$, i.e. $E_i \cap E_j = \emptyset$ when $i \neq j$, then

$$P\left(\bigcup_{i=1}^{\infty} E_i\right) = \sum_{i=1}^{\infty} P(E_i)$$

Properties

- **Proposition:** $P(E^c) = 1 - P(E)$.

Properties

- **Proposition:** $P(E^c) = 1 - P(E)$.
- **Proposition:** If $E \subset F$ then $P(E) \leq P(F)$.

Properties

- **Proposition:** $P(E^c) = 1 - P(E)$.
- **Proposition:** If $E \subset F$ then $P(E) \leq P(F)$.
- **Proposition:** We have $P(E \cup F) = P(E) + P(F) - P(E \cap F)$.

Conditional Probabilities

- **Conditional Probability.** Consider an experiment with sample space S . Let E and F be two events, then the conditional probability of E given F is denoted by $P(E|F)$ and satisfies if $P(F) > 0$

$$P(E|F) = \frac{P(E \cap F)}{P(F)}$$

Conditional Probabilities

- **Conditional Probability.** Consider an experiment with sample space S . Let E and F be two events, then the conditional probability of E given F is denoted by $P(E|F)$ and satisfies if $P(F) > 0$

$$P(E|F) = \frac{P(E \cap F)}{P(F)}$$

- **Intuition:** If F has occurred, then, in order for E to occur, it is necessary that the occurrence be both in E and F , hence it must be in $E \cap F$. Once F has occurred, F is the new sample space.

Conditional Probabilities

- *Equally likely outcomes.* In this case, we have

$$\begin{aligned} P(E|F) &= \frac{\# \text{ outcomes in } E \cap F}{\# \text{ outcomes in } F} \\ &= \underbrace{\frac{\# \text{ outcomes in } E \cap F}{\# \text{ outcomes in } S}}_{P(E \cap F)} / \underbrace{\left(\frac{\# \text{ outcomes in } F}{\# \text{ outcomes in } S} \right)}_{P(F)}. \end{aligned}$$

Independence

- Events A and B are independent iff $P(A \cap B) = P(A)P(B)$

Independence

- Events A and B are independent iff $P(A \cap B) = P(A)P(B)$
- Equivalent to $P(A|B) = P(A)$

Independence

- Events A and B are independent iff $P(A \cap B) = P(A)P(B)$
- Equivalent to $P(A|B) = P(A)$
- One event occurring does not effect the probability of another occurring

The Multiplication Rule

Let E_1, E_2, \dots, E_n be a sequence of events, then we have

$$P(E_1 \cap E_2 \cap \dots \cap E_n) = P(E_1) P(E_2 | E_1) \times \\ \times P(E_3 | E_1 \cap E_2) \dots P(E_n | E_1 \cap \dots \cap E_{n-1})$$

Example

- **Example:** You have a box with 3 blue marbles, 2 red marbles, and 4 yellow marbles. You are going to pull out one marble, record its color, put it back in the box and draw another marble. What is the probability of pulling out a red followed by a blue?

Example

- **Example:** You have a box with 3 blue marbles, 2 red marbles, and 4 yellow marbles. You are going to pull out one marble, record its color, put it back in the box and draw another marble. What is the probability of pulling out a red followed by a blue?
- **Example:** Consider the same box of marbles. However, we are going to pull out the first marble, leave it out and then pull out the second marble. What is the probability of pulling out a red marble followed by a blue marble?

Random Variables

- A **random variable** is a function $R : S \rightarrow R$

Random Variables

- A **random variable** is a function $R : S \rightarrow R$
- Domain of R is the sample space S

Random Variables

- A **random variable** is a function $R : S \rightarrow R$
- Domain of R is the sample space S
- Range of R is the real line

Random Variables

Example: **Discrete Random Variable**

Experiment: flip 10 coins

Desired outcome: the number of heads

We care about: the number of heads that appear among 10 tosses (not the probability of getting a particular sequence of heads and tails)

Random Variables

Example: **Discrete Random Variable**

Experiment: flip 10 coins

Desired outcome: the number of heads

We care about: the number of heads that appear among 10 tosses (not the probability of getting a particular sequence of heads and tails)

Probability of a random variable R taking on some specific value k is:
 $P(R = k) = P(\{s : R(s) = k\})$, with $R(s)$ - number of heads occurring after s tosses

Random Variables

Example: **Continuous Random Variable**

$R(s)$ - random variable indicating the amount of time it takes for a fast food burger to decay

Random Variables

Example: **Continuous Random Variable**

$R(s)$ - random variable indicating the amount of time it takes for a fast food burger to decay

Probability that R takes on a value between two real constants a and b is:

$$P(a \leq R \leq b) = P(\{s : a \leq R(s) \leq b\})$$

Probability Distribution

A probability distribution is a summary of probabilities for the values of a random variable.

– It is a list/table/equation that links all possible outcomes of a random variable to their corresponding probability values.

Probability Distribution

A probability distribution is a summary of probabilities for the values of a random variable.

– It is a list/table/equation that links all possible outcomes of a random variable to their corresponding probability values.

- **Mean** is the arithmetical average value of the data.
- **Median** is the middle value of the data.
- **Mode** is the most frequently occurring value of the data.
- **Expected value** of some a random variable X with respect to a distribution $P(X=x)$ is the mean value of X when x is drawn from P .
- **Variance** is the measure of variability in the data from the mean value.

Probability Distribution

Binomial: the random variable can have only two outcomes.

```
import numpy as np
n=100 # number of trials
p=0.5 # probability of success
s=1000 # size
np.random.binomial(n,p,s)
```

Probability Distribution

Binomial: the random variable can have only two outcomes.


```
import numpy as np
n=100 # number of trials
p=0.5 # probability of success
s=1000 # size
np.random.binomial(n,p,s)
```

Uniform: equal likelihood.

```
import numpy as np
np.random.uniform(low=1, high=10,size=100)
```


Probability Distribution

Normal (Gaussian): most common.



Bayes's theorem

Bayesian approach provides mathematical rule explaining how you should change your existing beliefs in the light of new evidence.

Bayes's theorem

- $posterior = \frac{likelihood * prior}{marginal\ likelihood}$

Bayes's theorem

- $posterior = \frac{likelihood * prior}{marginal\ likelihood}$
- $P(R = r|e) = \frac{P(e|R=r)P(R=r)}{P(e)}$
- $P(R = r|e)$: probability that random variable R has value r given evidence e

Bayes's theorem

- $\text{posterior} = \frac{\text{likelihood} * \text{prior}}{\text{marginal likelihood}}$
- $P(R = r|e) = \frac{P(e|R=r)P(R=r)}{P(e)}$
- $P(R = r|e)$: probability that random variable R has value r given evidence e
- The denominator is just a normalizing constant (called *marginal likelihood*) that ensures the posterior adds up to 1; it can be computed by summing up the numerator over all possible values of R , i.e.,
$$P(e) = P(R = 0, e) + P(R = 1, e) + \dots = \sum_r P(e|R = r)P(R = r)$$

Naive Bayes Algorithm

- Simple (“naive”) classification method based on Bayes rule
- Relies on very simple representation of document:
 - e.g. “bag of words”

Text Classification

Input:

- document d
- fixed set of classes $C = \{c_1, c_2, \dots, c_j\}$

Output: predicted class $c \in C$

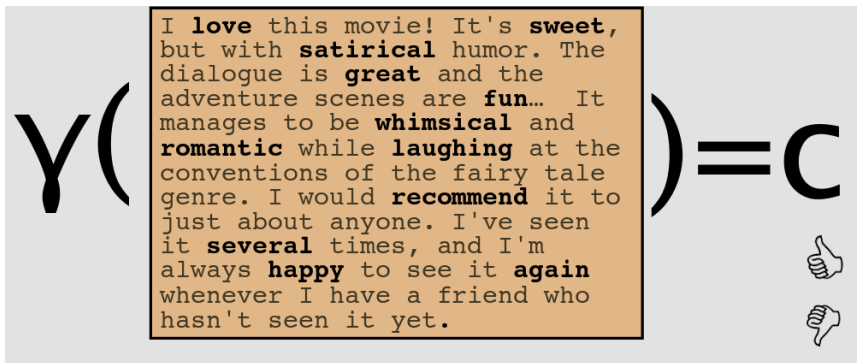
Supervised Learning for Text Classification

Input:

- document d
- fixed set of classes $C = \{c_1, c_2, \dots, c_j\}$
- A training set of m hand-labeled documents $(d_1, c_1), \dots, (d_m, c_m)$

Output: a learned classifier $\gamma : d \rightarrow c$

Naive Bayes Algorithm



Naive Bayes Algorithm

$$Y(\text{X love XXXXXXXXXXXXXXXXXXXX sweet
XXXXXXXXX satirical XXXXXXXXXXXX
XXXXXXXXXXXXX great XXXXXXXX
XXXXXXXXXXXXXXXXXXXXXXXXX fun XXXX
XXXXXXXXXXXXXXXXX whimsical XXXX
romantic XXXX laughing
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
XXXXXXXXXXXXXXXXX recommend XXXXX
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
XX several XXXXXXXXXXXXXXXXXXXX
XXXXX happy XXXXXXXX again
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
XXXXXXXXXXXXXXXXXXXX}) = C$$

thumbs up
thumbs down

Naive Bayes Algorithm

$Y(\text{Table}) = C$

great	2
love	2
recommend	1
laugh	1
happy	1
...	...

👍
👎

Naive Bayes Algorithm

For a document d and a class c

$$P(c|d) = \frac{P(d|c)P(c)}{P(d)}$$

Naive Bayes Algorithm

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(c | d)$$

MAP is “maximum a posteriori” = most likely class

$$= \operatorname{argmax}_{c \in C} \frac{P(d | c)P(c)}{P(d)}$$

Bayes Rule

$$= \operatorname{argmax}_{c \in C} P(d | c)P(c)$$

Dropping the denominator

Naive Bayes Algorithm

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(d | c) P(c)$$

$$= \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n | c) P(c)$$

Document d
represented as
features
 $x_1 \dots x_n$

Naive Bayes Algorithm

$$c_{MAP} = \operatorname{argmax}_{c \in C} P(x_1, x_2, \dots, x_n | c) P(c)$$

$O(|X|^n \cdot |C|)$ parameters

Could only be estimated if a very, very large number of training examples was available.

How often does this class occur?

We can just count the relative frequencies in a corpus

Binarized (Boolean feature) Multinomial Naive Bayes

Intuition:

- Word occurrence may matter more than word frequency
- The occurrence of the word *fantastic* tells us a lot
- The fact that it occurs 5 times may not tell us much more

Binarized (Boolean feature) Multinomial Naive Bayes

Intuition:

- Word occurrence may matter more than word frequency
- The occurrence of the word *fantastic* tells us a lot
- The fact that it occurs 5 times may not tell us much more

Boolean Multinomial Naive Bayes

Clips all the word counts in each document at 1

Multinomial Naive Bayes

Assumption:

- Bag of Words: assume position does not matter.
- Conditional Independence: Assume the feature probabilities $P(x_i|c_j)$ are independent given the class c .

$$P(x_1, \dots, x_n|c) = P(x_1|c) \cdot P(x_2|c) \cdots P(x_n|c)$$

Multinomial Naive Bayes

Assumption:

- Bag of Words: assume position does not matter.
- Conditional Independence: Assume the feature probabilities $P(x_i|c_j)$ are independent given the class c .

$$P(x_1, \dots, x_n|c) = P(x_1|c) \cdot P(x_2|c) \cdots P(x_n|c)$$

Applying Multinomial Naive Bayes Classifiers to Text Classification

- features are generated from a simple multinomial distribution.
- multinomial distribution: probability of observing counts among a number of categories.