

# Classification scikit-learn

Artificial Intelligence @ Allegheny College

Janyl Jumadinova

September 29–October 1, 2021

# scikit-learn

- Popular Python machine learning library
- Designed to be a well documented and approachable for non-specialist
- Built on top of NumPy and SciPy
- `scikit-learn` can be easily installed with `pip` or `conda`

```
pip install scikit-learn
```

```
conda install scikit-learn
```

# Data representation in scikit-learn

- Training dataset is described by a pair of matrices, one for the input data and one for the output.
- Most commonly used data formats are a NumPy `ndarray` or a Pandas `DataFrame` / `Series`.

# Data representation in scikit-learn

- Training dataset is described by a pair of matrices, one for the input data and one for the output.
- Most commonly used data formats are a NumPy `ndarray` or a Pandas `DataFrame` / `Series`.
- Each row of these matrices corresponds to one sample of the dataset.
- Each column represents a quantitative piece of information that is used to describe each sample (called “features”).

# Data representation in scikit-learn

Feature Matrix ( $X$ )

$n\_features \rightarrow$

$\leftarrow n\_samples$


Target Vector ( $y$ )

$\leftarrow n\_samples$


image credit: James Bourbeau

# Features in scikit-learn

feature **Module**

<https://scikit-image.org/docs/dev/api/skimage.feature.html>

# Local Binary Pattern Feature Extraction

Introduced by Ojala et. al in “Multiresolution Gray Scale and Rotation Invariant Texture Classification with Local Binary Patterns”

- 1 Check whether the points surrounding the central point are greater than or less than the central point → get LBP codes (stored as array).
- 2 Calculate a histogram of LBP codes as a feature vector.




image credit:

[https://scikit-image.org/docs/dev/auto\\_examples/features\\_detection/plot\\_local\\_binary\\_pattern.html](https://scikit-image.org/docs/dev/auto_examples/features_detection/plot_local_binary_pattern.html)

# Local Binary Pattern Feature Extraction

- Example: The histogram of the LBP outcome is used as a measure to classify textures.




image credit:



# Estimators in scikit-learn

- Algorithms are implemented as **estimator** classes in `scikit-learn`.
- Each estimator in `scikit-learn` is extensively documented (e.g. the `KNeighborsClassifier` documentation) with API documentation, user guides, and example usages.
- A **model** is an instance of one of these estimator classes

# Training a model

## fit then predict

```
# Fit the model  
model.fit(X, y)
```

```
# Get model predictions  
y_pred = model.predict(X)
```

# Decision Tree in scikit-learn

```
from sklearn.tree import DecisionTreeClassifier
```

```
clf = DecisionTreeClassifier(max_depth=2)
```

```
clf.fit(X, y)
```

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=2,  
                        max_features=None, max_leaf_nodes=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, presort=False, random_state=None,  
                        splitter='best')
```

image credit: James Bourbeau

# Model performance metrics

- Many commonly used performance metrics are built into the metrics subpackage in scikit-learn.
- However, a user-defined scoring function can be created using the `sklearn.metrics.make_scorer` function.

```
# Classification metrics
from sklearn.metrics import (accuracy_score, precision_score,
                             recall_score, f1_score, log_loss)

# Regression metrics
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
y_pred = [0, 2, 1, 3, 1]
y_true = [0, 1, 1, 3, 2]
```

```
accuracy_score(y_true, y_pred)
```

0.6

```
mean_squared_error(y_true, y_pred)
```

0.4

# Separate training and testing sets

- scikit-learn has a convenient `train_test_split` function that randomly splits a dataset into a testing and training set.

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=2)

print(f'X.shape = {X.shape}')
print(f'X_test.shape = {X_test.shape}')
print(f'X_train.shape = {X_train.shape}')

X.shape = (150, 2)
X_test.shape = (30, 2)
X_train.shape = (120, 2)
```


image credit: James Bourbeau

# Model selection - hyperparameter optimization


- Model **hyperparameter** values (parameters whose values are set before the learning process begins) can be used to avoid under- and over-fitting.

# Model selection - hyperparameter optimization


- Model **hyperparameter** values (parameters whose values are set before the learning process begins) can be used to avoid under- and over-fitting.
- **Under-fitting** - model isn't sufficiently complex enough to properly model the dataset at hand.
- **Over-fitting** - model is too complex and begins to learn the noise in the training dataset.



Janyl Jumadinova



Classification scikit-learn



September 29–October 1, 2021

13 / 18

# k-fold cross validation

- **Cross-validation** is a resampling procedure used to evaluate machine learning models on a limited data sample.
- It uses a limited sample in order to estimate how the model is expected to perform in general when used to make predictions on data not used during the training of the model.
- The parameter **k** refers to the number of groups that a given data sample is to be split into.



# k-fold cross validation

1. Shuffle the dataset randomly.
2. Split the dataset into  $k$  groups.
3. For each unique group:
  - 3.1. Take the group as a hold out or test data set.
  - 3.2. Take the remaining groups as a training data set.
  - 3.3. Fit a model on the training set and evaluate it on the test set.
  - 3.4. Retain the evaluation score and discard the model.
4. Summarize the skill of the model using the sample of model evaluation scores.

# k-fold cross validation




image credit:

[https://scikit-learn.org/stable/modules/cross\\_validation.html](https://scikit-learn.org/stable/modules/cross_validation.html)

# k-fold cross validation



Image source: Raschka, Sebastian, and Vahid Mirjalili. [Python Machine Learning](#), 2nd Ed. Packt Publishing, 2017.

image credit:

James Bourbeau

# Cross Validation in scikit-learn

```
from sklearn.model_selection import cross_validate

clf = DecisionTreeClassifier(max_depth=2)
scores = cross_validate(clf, X_train, y_train,
                        scoring='accuracy', cv=10,
                        return_train_score=True)

print(scores.keys())
test_scores = scores['test_score']
train_scores = scores['train_score']
print(test_scores)
print(train_scores)

print('\n10-fold CV scores:')
print(f'training score = {np.mean(train_scores)} +/- {np.std(train_scores)}')
print(f'validation score = {np.mean(test_scores)} +/- {np.std(test_scores)}')

dict_keys(['fit_time', 'score_time', 'test_score', 'train_score'])
[0.78571429 0.64285714 0.83333333 0.66666667 1.          0.91666667
 0.54545455 0.72727273 0.81818182 0.72727273]
[0.79245283 0.79245283 0.76851852 0.80555556 0.75          0.77777778
 0.79816514 0.79816514 0.78899083 0.79816514]

10-fold CV scores:
training score = 0.787024375076132 +/- 0.016054059411612778
validation score = 0.7663419913419914 +/- 0.12718955265834164
```

image credit:

James Bourbeau