Classification, Object Detection

Artificial Intelligence @ Allegheny College

Janyl Jumadinova

September 20–24, 2021

Classification Formalized

- Observations are classified into two or more classes, represented by a response variable Y taking values $1, 2, ..., K$.
- We have a feature vector $X = (X_1, X_2, ..., X_p)$, and we hope to build a classification rule $C(X)$ to assign a class label to an individual with feature X.
- We have a sample of pairs $(y_i, x_i), i = 1, ..., N$. Note that each of the x_i are vectors.

Object/Recognition

Goal:

Find an object of a pre-defined class in a static image or video frame.

Object/Recognition

Goal:

Find an object of a pre-defined class in a static image or video frame.

Approach:

- Extract certain image features, such as edges, color regions, textures, contours, etc.
- Use some heuristics to find configurations and/or combinations of those features specific to the object of interest.
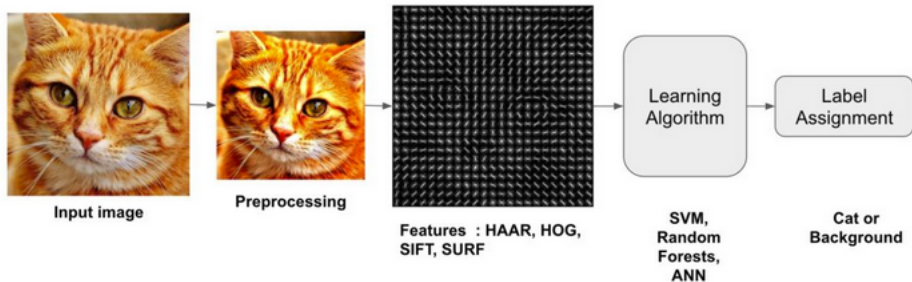
Statistical Model Training

- Training Set (Positive Samples/Negative Samples)
- Different features are extracted from the training samples and distinctive features that can be used to classify the object are selected.

Statistical Model Training

- Training Set (Positive Samples/Negative Samples)
- Different features are extracted from the training samples and distinctive features that can be used to classify the object are selected.
- Each time the trained classifier does not detect an object (misses the object) or mistakenly detects the absent object (gives a false alarm), model is adjusted.

Process of Object Detection/Recognition

Input image   Preprocessing   Features : HAAR, HOG, SIFT, SURF   Learning Algorithm   Label Assignment

SVM, Random Forests, ANN   Cat or Background

# Histogram of Oriented Gradients (HoG)

Histogram of Oriented Gradients (HoG)

From Dalal and Triggs paper

Linear classifiers

Find linear function to separate positive and negative examples

$\mathbf{x}_i$ positive : $\quad \mathbf{x}_i \cdot \mathbf{w} + b \geq 0$

$\mathbf{x}_i$ negative : $\quad \mathbf{x}_i \cdot \mathbf{w} + b < 0$

Which line
is best?

Support Vector Machines (SVMs)

- Discriminative classifier based on optimal separating line (for 2D case)
- Maximize the margin between the positive and negative training examples

Support Vector Machines (SVMs)

- Want line that maximizes the margin.

OpenCV: HOG and SVM for Person Detection

1. Represent each example with a single, fixed HoG template

2. Learn a single [linear] SVM as a detector

Decision Tree

Decision Tree

- **Root node**
  - Entry point to a collection of data
- **Inner nodes (among which the root node)**
  - A question is asked about data
  - One child node per possible answer
- **Leaf nodes**
  - Correspond to the decision to take (or conclusion to make) if reached

Decision Tree

- Represented by a series of binary splits.
- Each internal node represents a value query on one of the variables — e.g. "Is $X_3 > 0.4$". If the answer is "Yes", go right, else go left.

Decision Tree

- Represented by a series of binary splits.
- Each internal node represents a value query on one of the variables —
  e.g. "Is $X_3 > 0.4$". If the answer is "Yes", go right, else go left.
- The terminal nodes are the decision nodes.
- New observations are classified by passing their X down to a terminal
  node of the tree, and then using majority vote.

Decision Tree

✔ Can handle huge datasets

✔ Can handle mixed predictors—quantitative and qualitative

✔ Easily ignore redundant variables

✔ Handle missing data elegantly

✔ Small trees are easy to interpret

✘ large trees are hard to interpret

✘ Often prediction performance is poor

Decision Tree
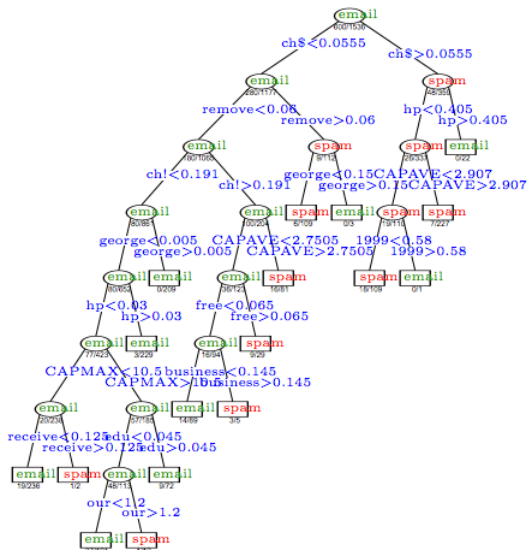
Model Averaging

Classification trees can be simple, but often produce noisy and weak classifiers.

- **Bagging**: Fit many large trees to bootstrap-resampled versions of the training data, and classify by majority vote.
- **Boosting**: Fit many large or small trees to reweighted versions of the training data. Classify by weighted majority vote.
- **Random Forests**: Fancier version of bagging.

Model Averaging

Classification trees can be simple, but often produce noisy and weak classifiers.

- **Bagging**: Fit many large trees to bootstrap-resampled versions of the training data, and classify by majority vote.
- **Boosting**: Fit many large or small trees to reweighted versions of the training data. Classify by weighted majority vote.
- **Random Forests**: Fancier version of bagging.

In general,

$$\text{Boosting} \succ \text{Random Forests} \succ \text{Bagging} \succ \text{Single Tree}.$$

Weak Classifier

- Computed feature value is used as input to a very simple decision tree classifier with 2 terminal nodes

$$\begin{cases} 1 & x_i \geq t_i \\ -1 & x_i \leq t_i \end{cases}$$

Boosted Classifier

- Complex and robust classifier is built out of multiple weak classifiers using a procedure called boosting.
- The boosted classifier is built iteratively as a weighted sum of weak classifiers.

Boosted Classifier

- Complex and robust classifier is built out of multiple weak classifiers using a procedure called boosting.
- The boosted classifier is built iteratively as a weighted sum of weak classifiers.
- On each iteration, a new weak classifier f_i is trained and added to the sum.

Boosted Classifier

- Complex and robust classifier is built out of multiple weak classifiers using a procedure called boosting.
- The boosted classifier is built iteratively as a weighted sum of weak classifiers.
- On each iteration, a new weak classifier f_i is trained and added to the sum.
- The smaller the error f_i gives on the training set, the larger is the coefficient/weight that is assigned to it.

Cascade of Boosted Classifiers

- Sequence of boosted classifiers with constantly increasing complexity.
- Chained into a cascade with the simpler classifiers going first.

OpenCV: Cascade Classifier

- Uses simple features and a cascade of boosted tree classifiers as a statistical model.
- Paul Viola and Michael J. Jones. " Rapid Object Detection using a Boosted Cascade of Simple Features." IEEE CVPR, 2001.

OpenCV: Cascade Classifier

- Classifier is trained on image of fixed size (Viola uses 24x24)
- Detection is done by sliding a search window of that size through the image and checking whether an image region at a certain location looks like our object or not.

OpenCV: Cascade Classifier

Feature's value is a weighted sum of two components:
- Pixel sum over the black rectangle
- Sum over the whole feature area

Cascade Classifier

- Instead of applying all the 6000 features on a window, group the features into different stages of classifiers and apply one-by-one.
- If a window fails the first stage, discard it. We don't consider remaining features on it.
- If it passes, apply the second stage of features and continue the process. The window which passes all stages is a face region.

OpenCV: Cascade Classifier

OpenCV already contains many pre-trained classifiers for face, eyes, smile etc. Those XML files are stored in `opencv/data/haarcascades/`
`cv2.CascadeClassifier.detectMultiScale(image[, scaleFactor[, minNeighbors[, flags[, minSize[, maxSize]]]]])`

- **scaleFactor** : Parameter specifying how much the image size is reduced at each image scale.
- **minNeighbors** : Parameter specifying how many neighbors each candidate rectangle should have to retain it. This parameter will affect the quality of the detected objects: higher value results in less detections but higher quality.

Classification Summary

- **Support Vector Machines (SVMs)**:
  - works for linearly separable and linearly inseparable data; works well in a highly dimensional space (text classification)
  - inefficient to train; probably not applicable to most industry scale applications

- **Random Forest**:
  - handle high dimensional spaces well, as well as the large number of training data; has been shown to outperform others

Classification Summary

No Free Lunch Theorem:

Wolpert (1996) showed that in a noise-free scenario where the loss function is the misclassification rate, if one is interested in off-training-set error, then there are no a priori distinctions between learning algorithms. On average, they are all equivalent.

Classification Summary

No Free Lunch Theorem:

Wolpert (1996) showed that in a noise-free scenario where the loss function is the misclassification rate, if one is interested in off-training-set error, then there are no a priori distinctions between learning algorithms. On average, they are all equivalent.

Occam's Razor principle:

Use the least complicated algorithm that can address your needs and only go for something more complicated if strictly necessary.

"Do we Need Hundreds of Classifiers to Solve Real World Classification Problems?"

http://jmlr.org/papers/volume15/delgado14a/delgado14a.pdf

OpenCV Cascade Training

<div align="center">

Training Cascade Tutorial

</div>

1. `opencv_createsamples`
2. `opencv_traincascade`
3. `opencv_visualisation`